

7CS108: Data Science and Data Mining

Anirban Chakraborty

a.chakraborty@wlv.ac.uk



Classification

Classification is a technique or a function that classifies given data into one of the pre-defined categories/classes.

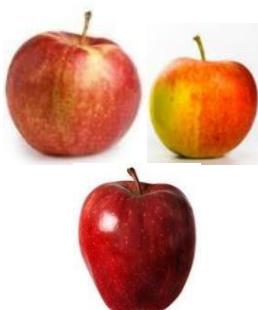


Apple



Pear

Distinguishing feature - colour





Classification

Classification is a technique or a function that classifies given data into one of the pre-defined categories/classes.

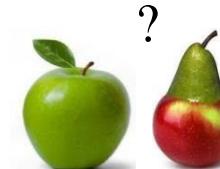
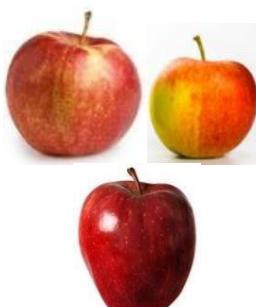


Apple



Pear

Distinguishing feature - colour



Is colour enough?

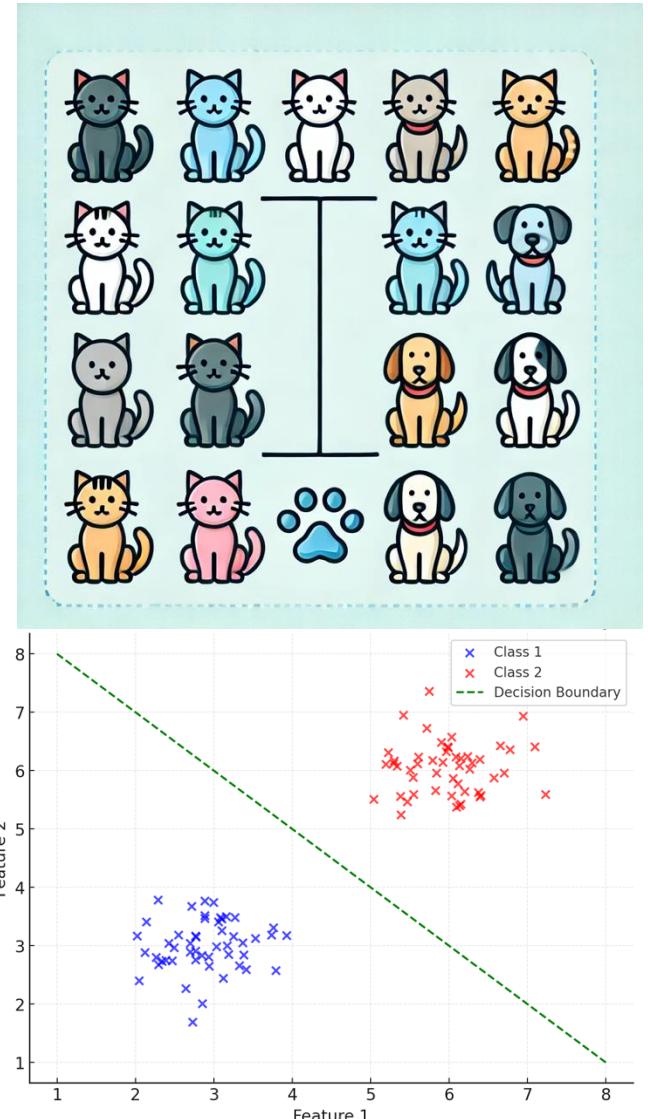
It may lead to misclassification!





Classification

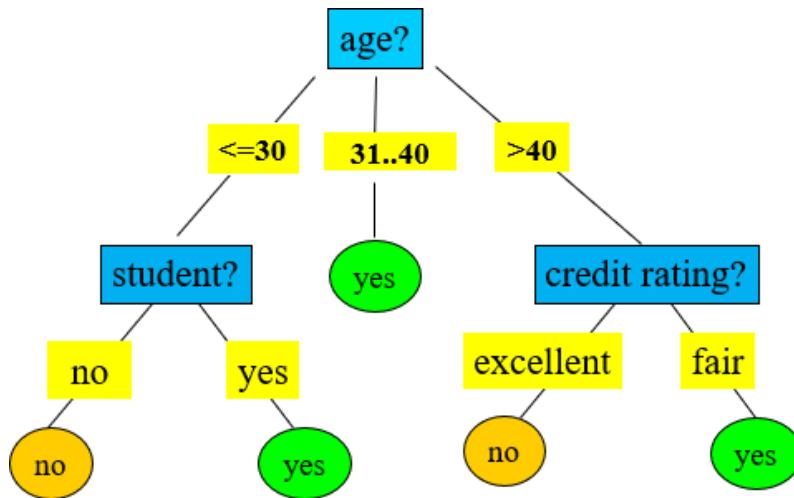
- Classification is a technique or a function that classifies given data into one of the pre-defined categories/classes.
- Classification is also known as “Supervised Learning”
 - There must be an “expert” (you) to “supervise” the computer.
 - In contrast, Clustering is known as “Unsupervised Learning”. We will discuss it in the later lectures.
 - There are semi-supervised and reinforcement learning too.
- From the data mining point of view...
 - Classification \approx Prediction \approx Forecasting
 - This is because the techniques are the same





Question?

Buys computer = yes ? no

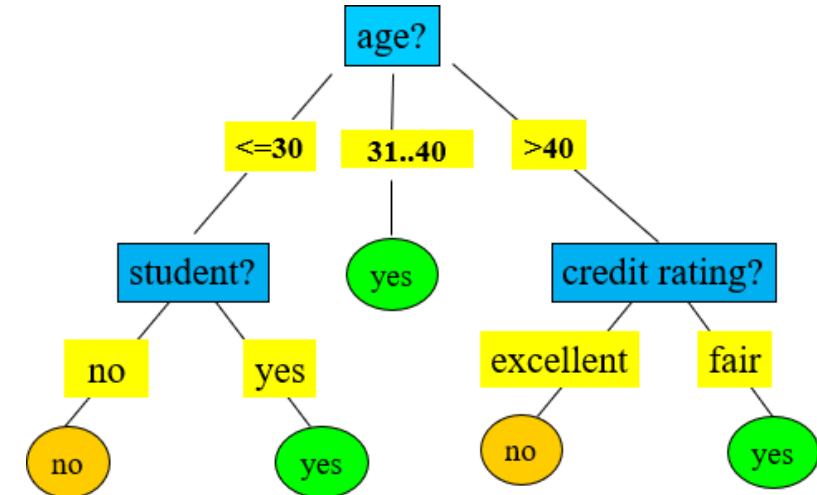


| Age | Student | Credit rating | Buys computer |
|-----|---------|---------------|---------------|
| 28 | Yes | Fair | ? |



What is Decision Tree?

- Flow-chart-like tree structure
 - Root node: a test on root node
 - Age < 30
 - Internal node: a test on an attribute
 - Student = yes ?
 - Leaf node: classes
 - Buys computer: yes or no?





Training set

- Whether a customer is likely to buy a computer?

| age | income | student | credit_rating | buys_computer |
|---------|--------|---------|---------------|---------------|
| <=30 | high | no | fair | no |
| <=30 | high | no | excellent | no |
| 31...40 | high | no | fair | yes |
| >40 | medium | no | fair | yes |
| >40 | low | yes | fair | yes |
| >40 | low | yes | excellent | no |
| 31...40 | low | yes | excellent | yes |
| <=30 | medium | no | fair | no |
| <=30 | low | yes | fair | yes |
| >40 | medium | yes | fair | yes |
| <=30 | medium | yes | excellent | yes |
| 31...40 | medium | no | excellent | yes |
| 31...40 | high | yes | fair | yes |
| >40 | medium | no | excellent | no |



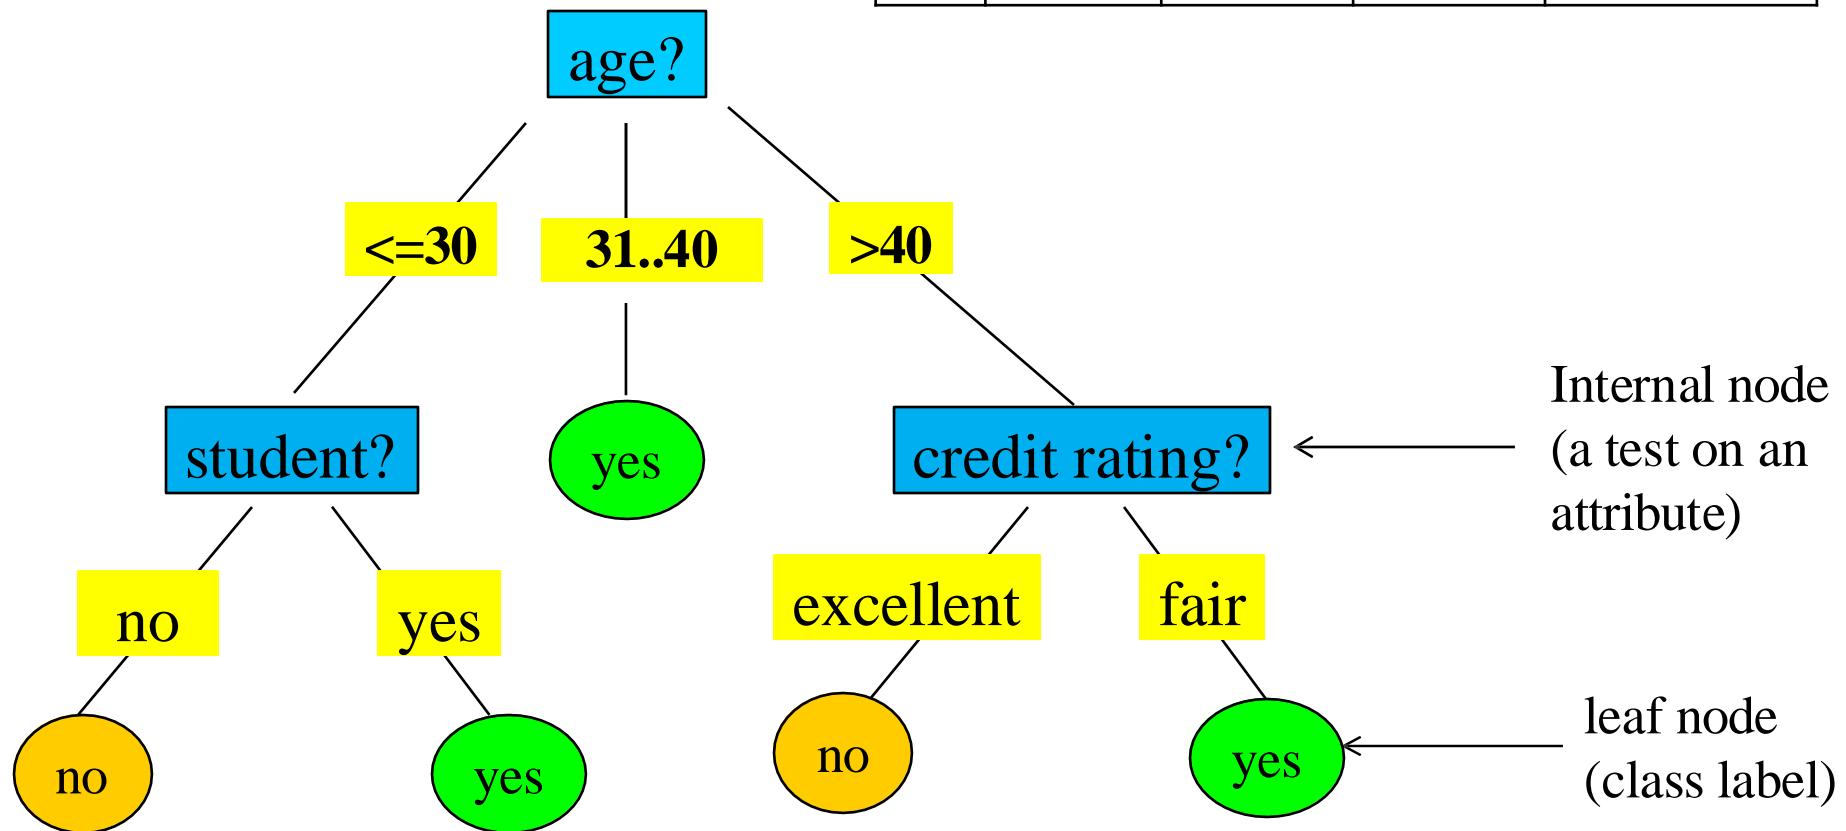
Data Preparation

- Data Cleaning
 - Pre-process data in order to reduce noise and handle missing values
- Relevance analysis (feature selection)
 - Remove the irrelevant or redundant attributes
- Data transformation
 - Generalise and/or normalise data



Decision Tree: “Buys Computer”

| Age | Income | Student | Credit rating | Buyς computer |
|-----|--------|---------|---------------|---------------|
| 25 | Low | Student | Fair | ? |





Decision Tree Induction

- Many algorithms:
 - Hunt's Algorithm
 - ID3 – late 1970's, J. Ross Quinlan
 - C4.5 – successor of ID3
 - CART – similar approach as C4.5
- General Structure
 - Attribute selection by calculating entropy and Information Gain
 - Decision tree induction
 - Rule generation



A Simple Thought on Decision Tree

- A simple minded algorithm:
 - If Age = a and Income = b and Student = c and Credit Rating = d
 - then buys_computer = ?
- This structure will yield a branch for each row, not eliminating superfluous attributes, and is unnecessarily complex.
- The complete classification space for m attributes is of size: $\prod_{j=1}^m N_j$
(N_j is the number of values of attribute j)



The Problem Statement

- Given a set of data described by **a set of attributes** and **an outcome class**, the problem is to find a **minimum decision tree** that will classify the values of the class based on the values of given attributes.



The Problem Statement

- Given a set of data described by **a set of attributes** and **an outcome class**, the problem is to find a **minimum decision tree** that will classify the values of the class based on the values of given attributes.
- How do we select an attribute?
 - Amongst four attributes: **AGE, INCOME, STUDENT** and **CREDIT RATING**, **which attribute carries “more” information than others?**



Information Theory

- Information Entropy → uncertainty
- Information Gain → reduce entropy

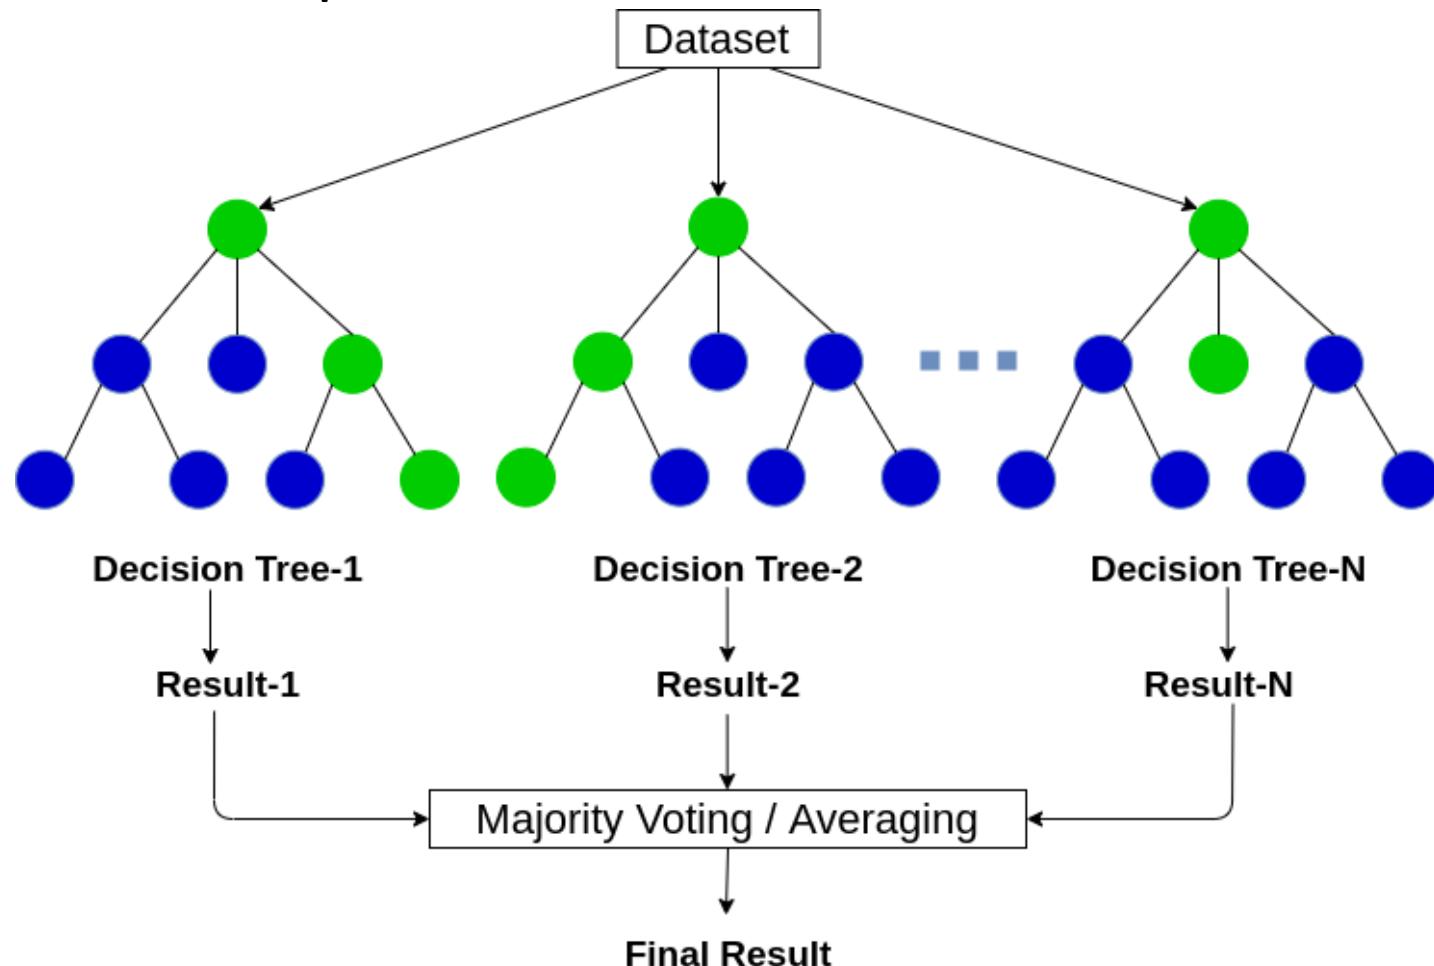
$$E = -\sum p_i \log_2(p_i)$$

$$\text{IG} = E(\text{Parent}) - \sum w_i E(\text{Child}_i)$$



Random Forest

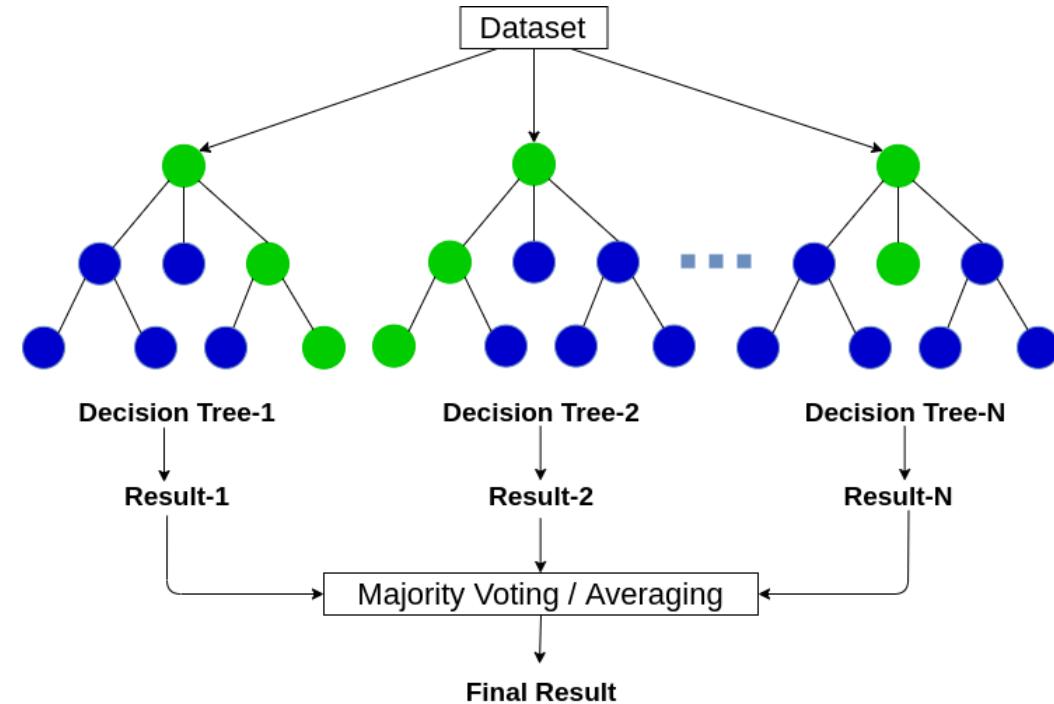
- Combines the output of multiple (randomly created) Decision Trees to generate the final output.





Random Forest

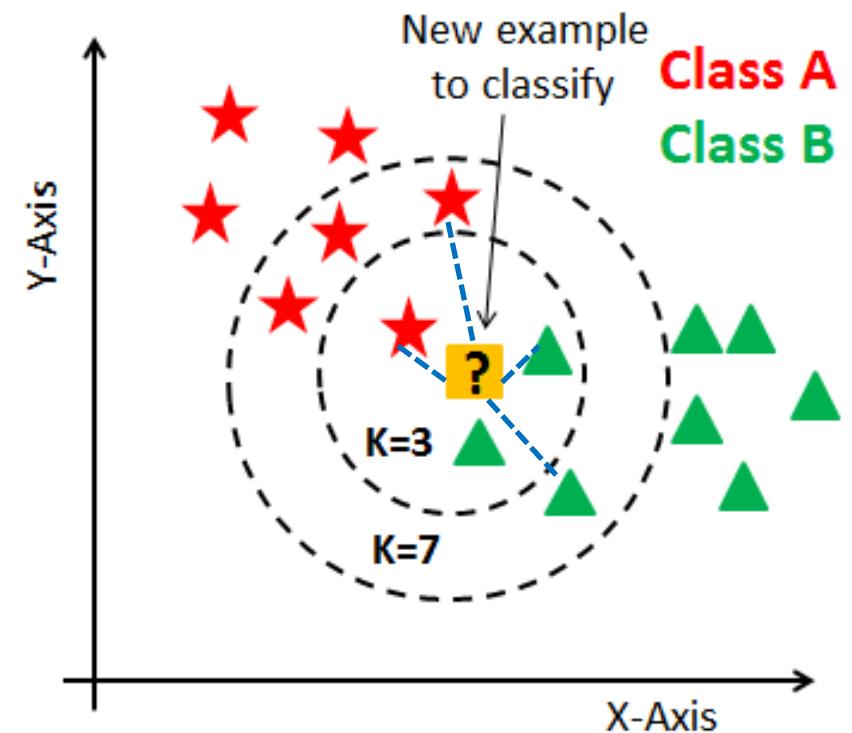
- **Individual Trees:** Each Decision Tree is trained on a random subset of the data (with replacement, i.e., bootstrapping) and a random subset of features.
- **Tree Predictions:** Each tree makes an independent prediction based on its learned splits and conditions.
- **Aggregation by Majority Voting:** The Random Forest combines the predictions of all trees, and the class with the highest votes becomes the final classification.





K-Nearest Neighbours (kNN)

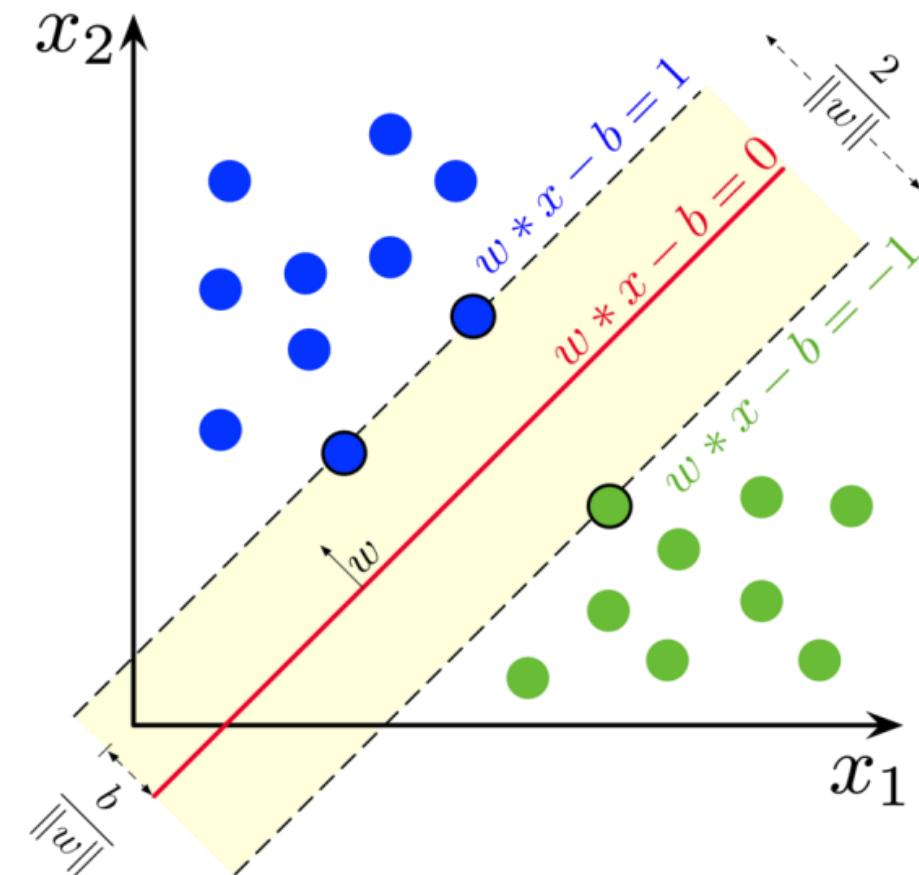
- Choose the value of k (number of neighbours).
- Compute the distance between the query point and all other data points.
- Identify the k nearest neighbours.
- Assign the class label based on the majority vote among k neighbours.
- In case of regression, take the mean of the k neighbours' values.





Support Vector Machine (SVM)

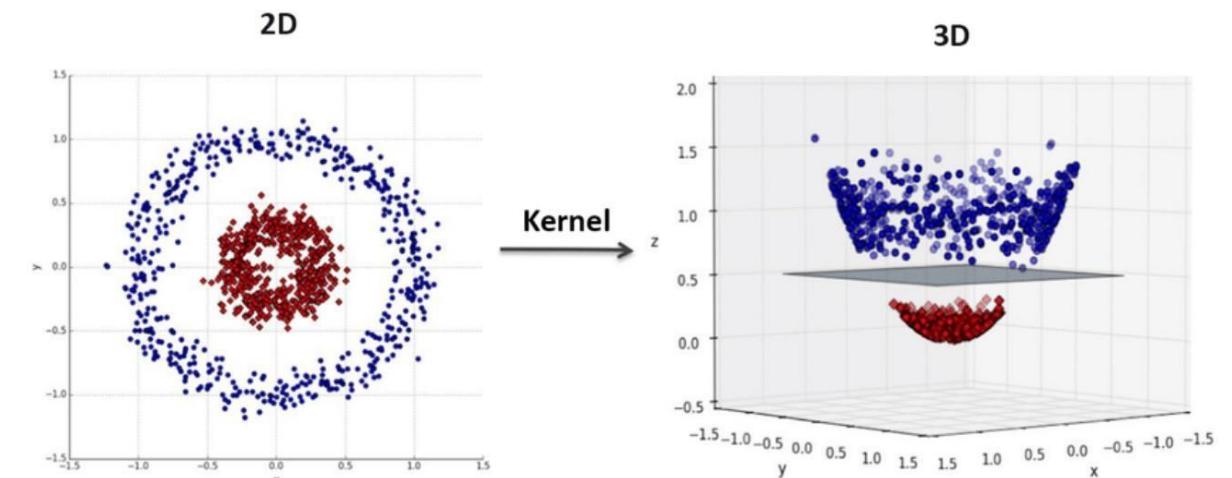
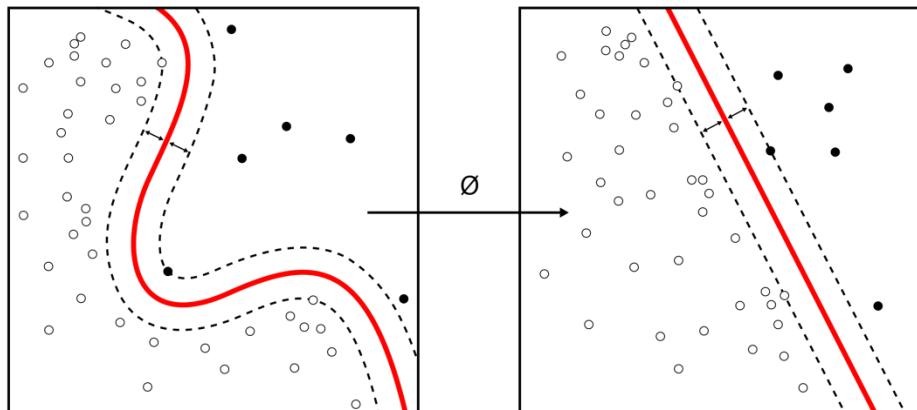
- The objective is to find a hyperplane in an n-dimensional space that separates the data points to their potential classes. The hyperplane should be positioned with the maximum distance to the data points.
- The data points with the minimum distance to the hyperplane are called Support Vectors





SVM

- The original maximum-margin hyperplane algorithm proposed by Vapnik in 1963 constructed a linear classifier that works when data are linearly separable (using a line or hyperplane)
- For data that are not linearly separable, we need to use a kernel function and transform data to another space where they are linearly separable. This is called the kernel trick.
- There are different kernel functions including linear function, Polynomial function, Radial basis function (RBF), and Sigmoid function.





Evaluation

- Confusion matrix
- Accuracy
- Precision
- Recall



Evaluation

Confusion Matrix

| | Actual Positive (Apple) | Actual Negative (Non-apple) |
|-----------------------------------|----------------------------|--------------------------------|
| Predicted Positive (Apple) | TP | FP |
| Predicted Negative (Non-apple) | FN | TN |

TP: true positive

FP: false positive

FN: false negative

TN: true negative

$$\text{Accuracy} = \frac{TP+TN}{TP+TN+FP+FN}$$

NOTE: Accuracy is NOT always a good measure!



Confusion Matrix

Total fruits: 20

| | Actual Positive (Apple) | Actual Negative (Non-apple) |
|-----------------------------------|----------------------------|--------------------------------|
| Predicted Positive (Apple) | 10 TP | 5 FP |
| Predicted Negative (Non-apple) | 2 FN | 3 TN |

TP: true positive

FN: false negative

FP: false positive

TN: true negative



Confusion Matrix

Total fruits: 20

| | Actual Positive (Apple) | Actual Negative (Non-apple) |
|-----------------------------------|----------------------------|--------------------------------|
| Predicted Positive (Apple) | 10 TP | 5 FP |
| Predicted Negative (Non-apple) | 2 FN | 3 TN |

True Positive (TP): the classifier correctly makes positive decisions

True Negative (TN): the classifier correctly makes negative decision

False Positive (FP): the classifier mistakenly makes positive decisions

False Negative (FN): the classifier mistakenly makes negative decisions



Limitation

Suppose

- Total number of fruits in the testing examples = 10,000
- Number of Non-apple = 9990
- Number of Apple = 10

Can you classify Apple?

If model predicts everything to be of class non-apple, the accuracy is $9990/10000 = 99.9\%!!!$

Here, accuracy is misleading because model cannot detect any Apple at all, still achieving high accuracy.



Evaluation

| | Actual Positive (Apple) | Actual Negative (Non-apple) |
|-----------------------------------|----------------------------|--------------------------------|
| Predicted Positive (Apple) | TP | FP |
| Predicted Negative (Non-apple) | FN | TN |

$$Precision = \frac{TP}{TP + FP}$$

$$Recall = \frac{TP}{TP+FN}$$

$$F_1 = 2 * \frac{Precision * Recall}{Precision + Recall}$$



Evaluation

| | Actual Positive (Apple) | Actual Negative (Non-apple) |
|-----------------------------------|----------------------------|--------------------------------|
| Predicted Positive (Apple) | TP | FP |
| Predicted Negative (Non-apple) | FN | TN |

$$Precision = \frac{TP}{TP + FP} = \frac{TP}{predicted\ positive}$$

Precision is the total number of **correctly identified actual apple cases out of retrieved apple**

$$Recall = \frac{TP}{TP+FN} = \frac{TP}{actual\ positive}$$

Recall is the number of **correctly identified apple cases from the total number of actual apple cases**



Evaluation

- Methods for Performance Evaluation
 - How to obtain reliable estimates?
 - i.e. how to partitioning the data?
- N-Fold cross validation
- Training and Testing sets



Cross Validation

- Ensures generalization to unseen data.
- Prevents overfitting.
- Provides reliable performance evaluation.
- Improves model robustness and reduces evaluation bias.
- Popular approaches:
 - K-Fold Cross Validation
 - Leave-One-Out





train_test_split()

```
x_train, x_test, y_train, y_test = train_test_split(features, targets, test_size=.30,  
random_state=50)
```

| features | | | | targets |
|----------|-------------|----------|--------|---------|
| age | competition | type | profit | targets |
| old | no | software | down | 0 |
| midlife | yes | software | up | 0 |
| midlife | no | hardware | down | 0 |
| old | no | hardware | down | 1 |
| young | no | hardware | up | 1 |
| young | no | software | up | 1 |
| midlife | no | software | up | 0 |
| young | yes | software | up | 1 |
| midlife | yes | hardware | down | 1 |
| old | yes | software | down | 1 |

Company.csv

